Trust of Artificial Intelligence in the Workplace: A Mixed-Methods Approach

Elizabeth Darnell
elizabeth.g.darnell@mytudublin.ie
Technological University Dublin
Dublin, Ireland

Abstract

Trust is critical for understanding individual use of generative artificial intelligence and as AI tools become ubiquitous in the workplace, it becomes more important to understand how individuals trust AI within that context. Research of trust of AI is relatively fragmented with significant reliance on custom scales. This research presents a mixed-methods protocol that evaluates the trust and distrust of AI within a workplace. The protocol utilises the Trust between People and Automation (TPA) scale which has been validated for use within AI contests and semi-structured interviews are used to help ground the research in the specific context being evaluated. Initial results of use of this protocol for evaluating trust of integrated development environments (IDE) indicate that genders trust AI differently. Further analysis is ongoing.

CCS Concepts

• **Human-centered computing → Human computer interaction (HCI)**; **User studies**.

Keywords

AI, Trust, Measurement

ACM Reference Format:
Elizabeth Darnell. 2026. Trust of Artificial Intelligence in the Workplace: A Mixed-Methods Approach. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26), April 13–17, 2026, Barcelona, Spain.* ACM, New York, NY, USA, 5 pages. https://doi.org/10.1145/3772363.3799220

1 Introduction

Generative artificial intelligence is being widely adopted and integrated into workplaces from healthcare to academia and technology and beyond [6]. The widespread adoption and deployment of AI tools into the workplace, despite productivity gains, is not without its challenges, including significant changes to the workforce and environmental concerns [4, 6, 18, 19]. In combination with both the adoption and the challenges, there is evidence of a strong public desire for the regulation of AI globally [6]. In the European Union, there is a greater importance being placed on developing and deploying trustworthy AI systems due to the EU AI Act [16, 20]. The EU AI Act is a risk-based framework that is structured around the importance of trustworthy AI [16, 20]. However, despite the significant focus on trustworthy AI, both within the EU and beyond, there has not been as much focus on how individuals trust AI which is a critical part to the discussion of and understanding of trustworthy AI. Additionally, there is research that indicates that trust is closely tied to reliance on tools which could help make better sense of adoption [1, 10, 11, 21, 32]. As such, it is important to improve the understanding and measurement of trust of AI in work contexts, especially as AI becomes more ubiquitous in the workplace. This research aims to develop a mixed-methods protocol to evaluate the trust of AI tools in the workplace and to improve the understanding of how individuals trust these tools as well as improve the methods of trust measurement.

1.1 Understanding and Measuring Trust

Understanding Trust. Researching how individuals trust is an inherently multidisciplinary venture which has lead to fractured methodologies and understandings [5, 14, 25, 28, 31]. The definition of trust that this research will use is "the attitude than an agent will help achieve an individual's goal in a situation characterized by uncertainty and vulnerability" [11, p. 54]. The important elements of this definition is that trust is an attitude, there exists "positive expectation", and that uncertainty and vulnerability or risk is present [11, 24, p. 395]. Different measurements approaches can be guided by the definition or conceptualisation of trust that is chosen [13]. However, this is not a consistent approach choosing methodologies and definitions further leading to the divergence in methods and understanding [13, 14].

One of the conflicts in both understanding and measuring trust is whether to conceptualise trust as a one-dimensional or two-dimensional concept [13, 25, 29]. An understanding of trust that is one-dimensional means that trust and distrust are on the opposite ends of the same continuum [13, 21, 25]. A two-dimensional understanding of trust is that trust and distrust are related but distinct concepts and that a high-level of trust does not necessarily indicate a low-level of distrust and vice versa [13, 21, 25]. An emerging convergence in the conceptualisations of trust, with measurement implications, is the two-dimensional understanding [13, 25, 29]. This research will utilise the two-dimensional understanding of trust.

Measuring Trust. Measuring trust is as complex as defining trust and there is significant diversity in methods to measure trust [5, 11, 13, 22, 25, 28]. Despite the diversity, there are some themes within these approaches particularly the use scales and qualitative methods [9, 11, 14, 28]. While the diverse methods to measure trust hinder reproducibility in researching trust, part of the reason for this landscape is that trust is a very context-specific concept which

This work is licensed under a Creative Commons Attribution 4.0 International License.
CHI EA '26, Barcelona, Spain
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2281-3/2026/04
https://doi.org/10.1145/3772363.3799220

leads researchers to rely on custom scales and methods to evaluate trust in unique contexts [9, 12, 14, 21].

Understanding and Measuring Trust of AI. Beyond understanding and measuring trust in a broader scope, there are significant questions related to how an individual might trust technology, especially generative AI because of the different understandings of trust that depend upon who or what the trustee is. Three of the primary types of trust are interpersonal trust, organisational trust, and trust of technology [11, 13, 17, 23]. Generative AI provides a complication in understanding which framework of trust should be utilised [8, 21]. The way in which generative AI tools are used and deployed complicates the end-users understanding of the tool itself, particularly due to the frequent anthropomorphisation of generative AI tools by the user [8, 30, 32]. Anthropomorphisation can help increase the perceived trustworthiness and perceived reliability of a tool [8]. The anthropomorphisation of these tools has the potential to shift the type of trust from trust of technology or organisational trust to interpersonal trust [14, 21]. It is unknown how the relationship between trusting the developers of the tool and trusting the tool or technology itself is related in overall trust of AI [8]. Trust of a specific technology and organisational trust remain important in understanding how end-users trust generative AI tools because a potential factor of trusting a generative AI tool is trusting the developers of the tool itself which is aligned with organisational trust [8].

One of the frameworks of trust follows the understanding that perceived trustworthiness leads to trust [8, 15]. A common framework of perceived trustworthiness is ability, benevolence, and integrity (ABI) [15]. However, it is unknown if objects or technology can have benevolence and integrity or if benevolence and integrity in the circumstance of trust of technology is related to the developer of the system [8].

While there are divergences within trust of technology research, there is an increased focus on understanding trust of AI in research and as such, Lockey and Gillespie [14] have proposed a research agenda with six main empirical gaps. The gaps are:

(1) Advancing a multi-stakeholder and multilevel approach to understanding trust in socio-technical systems,
(2) Understanding "optimal" trust in AI systems,
(3) Understanding AI trust failures and repair,
(4) Closing the gap between the principles and practice of trustworthy AI,
(5) Examining the AI trust intention-behaviour gap, and
(6) Enhancing the rigour and context-sensitivity of AI trust research and measurement [14, p. 13,15-18].

My research aims to make contributions to the empirical gaps 1, 5, and 6. Gap 1 is addressed through the qualitative portion of the following protocol with significant attention paid to the the context of the individual using AI tools within their workplace. Gap 5 is addressed via the behavioural and attitudinal measures and investigations of the following proposed protocol. Gap 6 is addressed through the following protocol that uses validated measures as well as qualitative components to improve both the rigour of trust research while maintaining rich grounding in the context being examined.

2 Research

My research aims to develop and test a mixed-methods protocol to evaluate trust of AI tools in the workplace and improve the understanding in how individuals trust AI tools. A mixed-methods approach has been chosen to ensure that the context of the trust environment is thoroughly explored as well as the use of a reliable scale to provide a comparable quantified measure. The research itself will be a series of three to five user studies with a variety of generative AI tools as the tested tool within a workplace context. The mixed-method protocol will be used in all studies and iterated upon with the ultimate aim for the protocol to be utilised by other researchers in other research contexts.

The studies will take place within the island of Ireland in a variety of professional environments. The participant selection requirements will differ related to each professional work environment that we conduct user studies within and participants will be recruited via snowball sampling.

2.1 Research Questions

The research questions are:

R1: How does the accuracy of the generative AI tools impact overall trust of AI?

R2: How does the self-described expertise in a specific area impact trust of AI?

R3: Does the personal demographics and traits of the participant impact trust of AI?

2.2 Protocol

The participants fill out an initial survey that covers their existing use of generative AI tools in their workplace that covers the type of tasks they use the tools for, the tools they use, and the frequency in which they use them. Next, participants complete a short task using a generative AI tool. The task will range from structured tasks to open-ended exploration. Following the task, they answer a short set of questions related to their confidence in the generated output of the tool and then complete the Trust between People and Automation (TPA) scale [7]. The final phase of the protocol is a semi-structured interview about the participant's experience with the task and tested tool and their general use of AI in the workplace. The protocol will be iterated upon after each study including the refinement of semi-structured interview questions and validation that the TPA scale remains reliable in new contexts.

2.3 Methods

The protocol utilises a mixed-methods approach to understand and investigate trust of AI in the workplace. For the quantitative portion of the protocol, the TPA scale is used [7, p. 30]. The TPA scale has 12-items and is measured using a seven-point Likert scale. This scale is a more commonly used scale in research of trust and technology, particularly in the new generative AI landscape [3, 14, 21, 25, 27]. There has been recent research indicates that the TPA scale is an appropriate scale to measure trust within the two-dimensional conceptualisation of trust despite its initial intention of a one-dimensional understanding [14, 21, 25, 26].

The qualitative data will be analysed using thematic analysis from

Braun and Clarke [2]. As trust is incredibly contextual, the incorporation of rich qualitative data and analysis allows for us increased context-awareness and depth to the understanding of how individuals conceptualise their trust or distrust of AI in the workplace.
This mixed methods approach allows for more rigorous evaluation of trust via the combination of the validated TPA scale as well as rich semi-structured interviews.

2.4 Study 1: Trust of Agentic IDEs in the Workplace

This initial study looked to evaluate trust of agentic integrated development environments (IDEs) in the workplace. As risk and uncertainty is required for the development of trust, we framed the task in the workplace context because there is risk and uncertainty in relying on generative AI tools in the workplace due to the impact on employment [11, 21, 24]. The task was 15 minutes of open-exploration with Windsurf, an agentic IDE. Participants had access to a sample prompt list to facilitate exploration if needed. Prompts ranged from basic data visualisation to developing a lightweight app that calculated nutritional content based on an ingredient list or analysing a generated data set. Most participants did not work off the prompt list and explored Windsurf using a coding task that they were familiar with from their work. These tasks ranged from physics problems, machine learning tasks, and hydraulic analyses. The data collection was completed in August and September of 2025. As of January 2026, data analysis is ongoing.

Participants. Participants were expert developers who had significant experience with programming. This study had 15 participants with more male representation (9 male and 6 female). The participants were primarily from academia in Dublin, Ireland (10 in academia, 4 in technology sector, and 1 in civil engineering). The participants were young with 66% being under 34 years old. All participants were residing and working in Ireland with 60% being from the EU.
This study was approved by the Research Ethics Committee at Technological University Dublin. All participants were provided and signed informed consent prior to data collection.

Research Questions. The research questions for this study follow the overall research questions as detailed above.

Methods. As detailed above, the qualitative data collected via the semi-structured interviews will analysed using thematic analysis following the Braun and Clarke method [2]. The quantitative results are analysed primarily following the structure utilised by Scharowski et al. in their analysis of the TPA scale using the two-dimensional understanding of trust [21, 25]. There was no change to the language of the TPA scale to improve the reliability and validity of the scale.

Initial Quantitative Results. We have completed the initial quantitative results from the TPA survey. These initial results have been obtained by using the statistical software R (version 4.5.1). To calculate the trust and distrust scores, the mean is taken. The trust items in the TPA are 6 - 12 and the distrust items are 1 - 5 with the distrust items being reversed. There was an average trust score of 3.38 and an average distrust score of 4.83. Summary statistics are reported in Table 1. Further analysis will be required but it is worth noting that individuals report higher levels of distrust than trust. This is particularly interesting given that 10 (66%) of the participants indicated that their AI use at work was $1-3$ hours daily and only 2 reported using AI at work for less than 3 hours a week. It will be worth exploring more on how frequent use of AI tools at work impact trust and distrust which is related to R2. To begin to answer R3, the difference in trust and distrust of AI

	Mean	Median	Standard Deviation
Trust Score	3.38	3.29	1.01
Distrust Score	4.83	4.4	1.56

Table 1: Summary statistics of trust and distrust scores

between the genders of the participants is investigated. The F-Test

	Trust	***SD*** Trust	Distrust	***SD*** Distrust
Women	2.64	0.8	4.33	1.60
Men	3.87	0.84	5.16	1.53

Table 2: Summary statistics of trust and distrust scores by gender

was completed to verify that the t-Test is appropriate for this dataset ($p_{trust} = 0.95, p_{distrust} = 0.86$). From the results in Table 3, there is evidence at $p < 0.05$ that there is a gender difference in trust levels of AI in this context. However, at $p < 0.05$ there is not evidence that there is a gender difference in distrust levels of AI in this context. Further analysis will be done to understand this difference in how

	t	df	p
Trust	2.83	13	0.01
Distrust	1.00	13	0.34

Table 3: T-test for gender differences

gender impacts the trust of AI within specific contexts.

Future Work. The analysis for Study 1 will be completed and further investigation into the gender differences of trust and distrust will be explored as well as the role of expertise and familiarity with tools that impact trust and distrust. Future research using this protocol will be informed by the first study via improvements to the semi-structured interview questions as well as context setting for the task that participants complete.
Future iterations of this study using the protocol include retrieval-augmented generated AI tools that will focus on workplace policy. We intend to do this RAGbot tool study using the School of Computer Science policy with participants from staff and faculty as well as researchers, including PhD and post-doctoral. The context of these RAGbot studies will differ with one study being an open-exploration of a RAGbot tool and one will have a specific question that the user needs to answer utilising the RAGbot tool. The purpose of this is to further investigate how self-described expertise impacts trust of AI tools.

We also intend to recruit beyond the very technical participant pool we have currently used, including potentially testing the RAG-bot outside of STEM faculty in the university setting. We plan to test this protocol with other AI tools in the workplace as time and discovery allows.

3 Discussion

3.1 Limitations

Since these studies focus on the workplace, the participants are not representative of the general public. Future studies with similar methods could help determine how members of the general public, especially with less technical knowledge and familiarity with AI, trust AI tools. Additionally, we evaluated the task within the concept of a work environment but it was not conducted in the workplace. Thus, we should be cautious in extrapolating directly how the participants trust AI in the workplace until further investigation and research is completed.

3.2 Contributions and Aims

My planned contributions to the body of knowledge is twofold: one is a validated protocol to evaluate trust of AI that can be applied in a variety of different contexts by other researchers and two is to contribute to the better understanding of how and why individual's trust AI in the workplace. These contributions are aligned with the research agenda proposed by Lockey and Gillespie [14]. In particular that the data collected via the semi-structured interviews aims to understand how individuals conceptualise AI tools in their broader environment and how their perceptions of the developers impact their trust. Additionally, I aim to contribute to the further understanding of how trust impacts behaviour through both qualitative and quantitative measures. The protocol and results from the protocol will be shared and published with the goal to improve the rigour of trust research while also maintaining context-sensitivity.

Declaration on Generative AI Use

The authors have utilised generative AI to debug and clarify code in quantitative analysis as well as provide minor phrase clarification and feedback. The authors primarily utilised Claude (Sonnet 4.5) by Anthropic.

Acknowledgments

I would like to thank and acknowledge my supervisors Dr. Dympna O'Sullivan and Dr. Emma Murphy for their support, feedback, and guidance in my research. I would also like to thank Dr. Elizabeth Resor for her feedback on this paper. This work was supported by the TU Dublin ARISE (Amplify Research & Innovation Supporting Enterprise) programme, funded under the TU RISE scheme, and co-financed by the Government of Ireland and the European Union through the ERDF Southern, Eastern & Midland Regional Programme 2021–2027 and the Northern & Western Regional Programme 2021–2027.

References

[1] Mathias Bollaert, Olivier Augereau, and Gilles Coppin. 2024. Measuring and Calibrating Trust in Artificial Intelligence. In *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, Vol. 14536 LNCS. Springer Science and Business Media Deutschland GmbH, York, United Kingdom, 232–237. doi:10.1007/978-3-031-61698-3{_}22

[2] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative Research in Psychology* 3, 2 (2006), 77–101. doi:10.1191/1478088706qp063oa

[3] Francesca Cabiddu, Ludovica Moi, Gerardo Patriotta, and David G. Allen. 2022. Why do users trust algorithms? A review and conceptualization of initial trust and trust over time. *European Management Journal* 40, 5 (10 2022), 685–706. doi:10.1016/j.emj.2022.06.001

[4] Olesya Dmitracova. 2025. 41% of companies worldwide plan to reduce workforces by 2030 due to AI. https://edition.cnn.com/2025/01/08/business/ai-job-losses-by-2030-intl/index.html

[5] Nicole Gillespie. 2012. Measuring trust in organizational contexts: an overview of survey-based measures. In *Handbook of Research Methods on Trust* (1 ed.), Fergus Lyon, Guido Möllering, and Mark N.K. Saunders (Eds.). Edward Elgar Publishing Limited, Cheltenham, UK, Chapter 17, 175–188.

[6] Nicole Gillespie, Steve Lockey, Tabi Ward, Alexandria Macdade, and Gerard Hassaed. 2025. *Trust, attitudes and use of artificial intelligence: A global study 2025.* Technical Report. The University of Melbourne and KPMG. 1–118 pages. doi:10.26188/28822919

[7] Jiun-Yin Jian, Ann M. Bisantz, and Colin G. Drury. 2000. Foundations for an Empirically Determined Scale of Trust in Automated Systems. *International Journal of Cognitive Ergonomics* 4, 1 (3 2000). doi:10.1207/s15327566ijce0401{_}04

[8] Kimberly M. Wingert and Roger C Mayer. 2024. Trust in autonomous technology: the technology or its maker? In *A Research Agenda for Trust: Interdisciplinary Perspectives*, Roger C. Mayer and Barbara M. Mayer (Eds.). Edward Elgar Publishing Limited, Cheltenham, UK, Chapter 5, 51–62.

[9] Moritz Körber. 2018. Theoretical considerations and development of a questionnaire to measure trust in automation. In *Proceedings 20th Triennial Congress of the IEA*. Springer, Florence, Italy, 1–20. https://github.com/moritzkoerber/TiA_Trust_in_Automation_Questionnaire

[10] Sebastian Krügel, Andreas Ostermaier, and Matthias Uhl. 2022. Zombies in the Loop? Humans Trust Untrustworthy AI-Advisors for Ethical Decisions. *Philosophy and Technology* 35, 17 (3 2022), 1–37. doi:10.1007/s13347-022-00511-9

[11] J. D. Lee and K. A. See. 2004. Trust in Automation: Designing for Appropriate Reliance. *Human Factors: The Journal of the Human Factors and Ergonomics Society* 46, 1 (2004), 50–80. doi:10.1518/hfes.46.1.50{_}30392

[12] Roy J. Lewicki and Chad Brinsfield. 2012. Measuring trust beliefs and behaviours. In *Handbook of Research Methods on Trust*, Fergus Lyon, Guido Möllering, and Mark Saunders (Eds.). Edward Elgar Publishing Limited, Cheltenham, Chapter 3, 29–39.

[13] Roy J. Lewicki, Edward C. Tomlinson, and Nicole Gillespie. 2006. Models of interpersonal trust development: Theoretical approaches, empirical evidence, and future directions. *Journal of Management* 32, 6 (12 2006), 991–1022. doi:10.1177/0149206306294405

[14] Steve Lockey and Nicole Gillespie. 2024. Understanding trust in artificial intelligence: a research agenda. In *A Research Agenda for Trust: Interdisciplinary Perspectives*, Roger C. Mayer and Barbara M. Mayer (Eds.). Edward Elgar Publishing, Cheltenham, UK, Chapter 2, 11–24.

[15] Roger C. Mayer, James H. Davis, and F. David Schoorman. 1995. An Integrative Model of Organizational Trust. *The Academy of Management Review* 20, 3 (7 1995), 709. doi:10.2307/258792

[16] Louise McCormack and Malika Bendechache. 2024. A Comprehensive Survey and Classification of Evaluation Criteria for Trustworthy Artificial Intelligence. *AI & Ethics* 5 (10 2024), 1973–1994. doi:10.1007/s43681-024-00590-8

[17] D. Harrison Mcknight, Michelle Carter, Jason Bennett Thatcher, and Paul F. Clay. 2011. Trust in a specific technology: An investigation of its components and measures. *ACM Transactions on Management Information Systems* 2, 2 (6 2011), 1–25. doi:10.1145/1985347.1985353

[18] Lex McMenamin. 2025. ChatGPT Is Everywhere — Why Aren't We Talking About Its Environmental Costs? https://www.teenvogue.com/story/chatgpt-is-everywhere-environmental-costs-oped

[19] Jordan Novet and Jonathan Vanian. 2025. Satya Nadella says as much as 30% of Microsoft code is written by AI. https://www.cnbc.com/2025/04/29/satya-nadella-says-as-much-as-30percent-of-microsoft-code-is-written-by-ai.html

[20] Publications Office of the European Union L and Luxembourg Luxembourg. 2024. Regulation (EU) 2024/1689 of the European Parliament and of the Council. http://data.europa.eu/eli/reg/2024/1689/oj

[21] Sebastian A.C. Perrig, Nicolas Scharowski, and Florian Brühlmann. 2023. Trust Issues with Trust Scales: Examining the Psychometric Quality of Trust Measures in the Context of AI. In *Conference on Human Factors in Computing Systems - Proceedings*. Association for Computing Machinery, Hamburg, Germany, 1–7. doi:10.1145/3544549.3585808

[22] Vlad L. Pop, Alex Shrewsbury, and Francis T. Durso. 2015. Individual differences in the calibration of trust in automation. *Human Factors* 57, 4 (6 2015), 545–556. doi:10.1177/0018720814564422

[23] Julian B Rotter. 1967. A new scale for the measurement of interpersonal trust. *Journal of Personality* 34, 4 (12 1967), 651–665. doi:10.1111/j.1467-6494.1967.

tb01454.x

[24] Denise M. Rousseau, Sim B. Sitkin, Ronald S. Burt, and Colin Camerer. 1998. Not So Different After All: A Cross-Discipline View Of Trust. *Academy of Management Review* 23, 3 (7 1998), 393–404. doi:10.5465/amr.1998.926617

[25] Nicolas Scharowski, Sebastian A. C. Perrig, Nick von Felten, Lena Fanya Aeschbach, Klaus Opwis, Philipp Wintersberger, and Florian Brühlmann. 2025. To Trust or Distrust AI: A Questionnaire Validation Study. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. ACM, New York, NY, USA, 361–374. doi:10.1145/3715275.3732025

[26] Randall D Spain, Ernesto A Bustamante, and James P Bliss. 2008. Towards an Empirically Developed Scale for System Trust: Take Two. *Proceedings of the Human Factors and Ergonomics Society Annual Meeting* 52, 19 (9 2008), 1335–1339. doi:10.1177/154193120805201907

[27] Takane Ueno, Yuto Sawa, Yeongdae Kim, Jacqueline Urakami, Hiroki Oura, and Katie Seaborn. 2022. Trust in Human-AI Interaction: Scoping Out Models, Measures, and Methods. In *Conference on Human Factors in Computing Systems - Proceedings*. Association for Computing Machinery, New Orleans, Louisiana, 1–7. doi:10.1145/3491101.3519772

[28] Oleksandra Vereschak, Gilles Bailly, and Baptiste Caramiaux. 2021. How to Evaluate Trust in AI-Assisted Decision Making? A Survey of Empirical Methodologies. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2 (10 2021), 1–39. doi:10.1145/3476068

[29] Roel Visser, Tobias M Peters, Ingrid Scharlau, Barbara Hammer, and Serge Thill. 2025. Trust, distrust, and appropriate reliance in (X)AI: A conceptual clarification of user trust and survey of its empirical evaluation. *Cognitive Systems Research* 91 (2025), 1–29. doi:10.1007/978-3

[30] Alan R. Wagner, Jason Borenstein, and Ayanna Howard. 2018. Overtrust in the robotic age. *Commun. ACM* 61, 9 (9 2018), 22–24. doi:10.1145/3241365

[31] Rongbin Yang and Santoso Wibowo. 2022. User trust in artificial intelligence: A comprehensive conceptual framework. *Electronic Markets* 32, 4 (12 2022), 2053–2077. doi:10.1007/s12525-022-00592-6

[32] Kun Yu, Shlomo Berkovsky, Ronnie Taib, Jianlong Zhou, and Fang Chen. 2019. Do I trust my machine teammate? An investigation from perception to decision. In *International Conference on Intelligent User Interfaces, Proceedings IUI*, Vol. Part F147615. Association for Computing Machinery, Marina del Ray, 460–468. doi:10.1145/3301275.3302277